

# Tutorial

This tutorial is intended as an introduction to working with **MongoDB** and **PyMongo**.

## Prerequisites

Before we start, make sure that you have the **PyMongo** distribution [installed](#). In the Python shell, the following should run without raising an exception:

```
>>> import pymongo
```

This tutorial also assumes that a MongoDB instance is running on the default host and port. Assuming you have downloaded and installed MongoDB, you can start it like so:

```
$ mongod
```

## Making a Connection with MongoClient

The first step when working with **PyMongo** is to create a `MongoClient` to the running **mongod** instance. Doing so is easy:

```
>>> from pymongo import MongoClient  
>>> client = MongoClient()
```

The above code will connect on the default host and port. We can also specify the host and port explicitly, as follows:

```
>>> client = MongoClient("localhost", 27017)
```

Or use the MongoDB URI format:

```
>>> client = MongoClient("mongodb://localhost:27017/")
```

## Getting a Database

A single instance of MongoDB can support multiple independent databases. When working with PyMongo you access databases using attribute style access on `MongoClient` instances:

```
>>> db = client.test_database
```

If your database name is such that using attribute style access won't work (like `test-database`), you can use dictionary style access instead:

```
>>> db = client["test-database"]
```

## Getting a Collection

A `collection` is a group of documents stored in MongoDB, and can be thought of as roughly the equivalent of a table in a relational database. Getting a collection in PyMongo works the same as getting a database:

```
>>> collection = db.test_collection
```

or (using dictionary style access):

```
>>> collection = db["test-collection"]
```

An important note about collections (and databases) in MongoDB is that they are created lazily - none of the above commands have actually performed any operations on the MongoDB server. Collections and databases are created when the first document is inserted into them.

## Documents

Data in MongoDB is represented (and stored) using JSON-style documents. In PyMongo we use dictionaries to represent documents. As an example, the following dictionary might be used to represent a blog post:

```
>>> import datetime
>>> post = {
...     "author": "Mike",
...     "text": "My first blog post!",
...     "tags": ["mongodb", "python", "pymongo"],
...     "date": datetime.datetime.now(tz=datetime.timezone.utc),
... }
```

Note that documents can contain native Python types (like `datetime.datetime` instances) which will be automatically converted to and from the appropriate BSON types.

# Inserting a Document

To insert a document into a collection we can use the `insert_one()` method:

```
>>> posts = db.posts  
>>> post_id = posts.insert_one(post).inserted_id  
>>> post_id  
ObjectId('...')
```

When a document is inserted a special key, `"_id"`, is automatically added if the document doesn't already contain an `"_id"` key. The value of `"_id"` must be unique across the collection. `insert_one()` returns an instance of `InsertOneResult`. For more information on `"_id"`, see the documentation on `_id`.

After inserting the first document, the `posts` collection has actually been created on the server. We can verify this by listing all of the collections in our database:

```
>>> db.list_collection_names()  
['posts']
```

# Getting a Single Document With `find_one()`

The most basic type of query that can be performed in MongoDB is `find_one()`. This method returns a single document matching a query (or `None` if there are no matches). It is useful when you know there is only one matching document, or are only interested in the first match. Here we use `find_one()` to get the first document from the `posts` collection:

```
>>> import pprint  
>>> pprint pprint(posts.find_one())  
{'_id': ObjectId('...'),  
 'author': 'Mike',  
 'date': datetime.datetime(...),  
 'tags': ['mongodb', 'python', 'pymongo'],  
 'text': 'My first blog post!'}
```

The result is a dictionary matching the one that we inserted previously.

## Note

The returned document contains an `"_id"`, which was automatically added on insert.

`find_one()` also supports querying on specific elements that the resulting document must match. To limit our results to a document with author "Mike" we do:

```
>>> pprint.pprint(posts.find_one({"author": "Mike"}))  
{'_id': ObjectId('...'),  
 'author': 'Mike',  
 'date': datetime.datetime(...),  
 'tags': ['mongodb', 'python', 'pymongo'],  
 'text': 'My first blog post!'}
```

If we try with a different author, like "Eliot", we'll get no result:

```
>>> posts.find_one({"author": "Eliot"})  
>>>
```

## Querying By ObjectId

We can also find a post by its `_id`, which in our example is an ObjectId:

```
>>> post_id  
ObjectId(...)  
>>> pprint.pprint(posts.find_one({"_id": post_id}))  
{'_id': ObjectId('...'),  
 'author': 'Mike',  
 'date': datetime.datetime(...),  
 'tags': ['mongodb', 'python', 'pymongo'],  
 'text': 'My first blog post!'}
```

Note that an ObjectId is not the same as its string representation:

```
>>> post_id_as_str = str(post_id)  
>>> posts.find_one({"_id": post_id_as_str}) # No result  
>>>
```

A common task in web applications is to get an ObjectId from the request URL and find the matching document. It's necessary in this case to **convert the ObjectId from a string** before passing it to `find_one`:

```
from bson.objectid import ObjectId

# The web framework gets post_id from the URL and passes it as a string
def get(post_id):
    # Convert from string to ObjectId:
    document = client.db.collection.find_one({'_id': ObjectId(post_id)})
```

#### See also

[When I query for a document by ObjectId in my web application I get no result](#)

## Bulk Inserts

In order to make querying a little more interesting, let's insert a few more documents. In addition to inserting a single document, we can also perform *bulk insert* operations, by passing a list as the first argument to `insert_many()`. This will insert each document in the list, sending only a single command to the server:

```
>>> new_posts = [
...     {
...         "author": "Mike",
...         "text": "Another post!",
...         "tags": ["bulk", "insert"],
...         "date": datetime.datetime(2009, 11, 12, 11, 14),
...     },
...     {
...         "author": "Eliot",
...         "title": "MongoDB is fun",
...         "text": "and pretty easy too!",
...         "date": datetime.datetime(2009, 11, 10, 10, 45),
...     },
... ]
>>> result = posts.insert_many(new_posts)
>>> result.inserted_ids
[ObjectId('...'), ObjectId('...')]
```

There are a couple of interesting things to note about this example:

- The result from `insert_many()` now returns two `ObjectId` instances, one for each inserted document.
- `new_posts[1]` has a different “shape” than the other posts - there is no `“tags”` field and we’ve added a new field, `“title”`. This is what we mean when we say that MongoDB is *schema-free*.

## Querying for More Than One Document

To get more than a single document as the result of a query we use the `find()` method. `find()` returns a `Cursor` instance, which allows us to iterate over all matching documents. For example, we can iterate over every document in the `posts` collection:

```
>>> for post in posts.find():
...     pprint.pprint(post)
...
{'_id': ObjectId('...'),
 'author': 'Mike',
 'date': datetime.datetime(...),
 'tags': ['mongodb', 'python', 'pymongo'],
 'text': 'My first blog post!'}
{'_id': ObjectId('...'),
 'author': 'Mike',
 'date': datetime.datetime(...),
 'tags': ['bulk', 'insert'],
 'text': 'Another post!'}
{'_id': ObjectId('...'),
 'author': 'Eliot',
 'date': datetime.datetime(...),
 'text': 'and pretty easy too!',
 'title': 'MongoDB is fun'}
```

Just like we did with `find_one()`, we can pass a document to `find()` to limit the returned results. Here, we get only those documents whose author is “Mike”:

```
>>> for post in posts.find({"author": "Mike"}):
...     pprint.pprint(post)
...
{'_id': ObjectId('...'),
 'author': 'Mike',
 'date': datetime.datetime(...),
 'tags': ['mongodb', 'python', 'pymongo'],
 'text': 'My first blog post!'}
{'_id': ObjectId('...'),
 'author': 'Mike',
 'date': datetime.datetime(...),
 'tags': ['bulk', 'insert'],
 'text': 'Another post!'}
```

## Counting

If we just want to know how many documents match a query we can perform a `count_documents()` operation instead of a full query. We can get a count of all of the documents in a collection:

```
>>> posts.count_documents({})
3
```

or just of those documents that match a specific query:

```
>>> posts.count_documents({"author": "Mike"})
2
```

## Range Queries

MongoDB supports many different types of advanced queries. As an example, lets perform a query where we limit results to posts older than a certain date, but also sort the results by author:

```
>>> d = datetime.datetime(2009, 11, 12, 12)
>>> for post in posts.find({"date": {"$lt": d}}).sort("author"):
...     pprint.pprint(post)
...
{'_id': ObjectId('...'),
 'author': 'Eliot',
 'date': datetime.datetime(...),
 'text': 'and pretty easy too!',
 'title': 'MongoDB is fun'}
{'_id': ObjectId('...'),
 'author': 'Mike',
 'date': datetime.datetime(...),
 'tags': ['bulk', 'insert'],
 'text': 'Another post!'}
```

Here we use the special `"$lt"` operator to do a range query, and also call `sort()` to sort the results by author.

## Indexing

Adding indexes can help accelerate certain queries and can also add additional functionality to querying and storing documents. In this example, we'll demonstrate how to create a [unique index](#) on a key that rejects documents whose value for that key already exists in the index.

First, we'll need to create the index:

```
>>> result = db.profiles.create_index([("user_id", pymongo.ASCENDING)], unique=True)
>>> sorted(list(db.profiles.index_information()))
['_id', 'user_id_1']
```

Notice that we have two indexes now: one is the index on `_id` that MongoDB creates automatically, and the other is the index on `user_id` we just created.

Now let's set up some user profiles:

```
>>> user_profiles = [{"user_id": 211, "name": "Luke"}, {"user_id": 212, "name": "Ziltoid"}]
>>> result = db.profiles.insert_many(user_profiles)
```

The index prevents us from inserting a document whose `user_id` is already in the collection:

```
>>> new_profile = {"user_id": 213, "name": "Drew"}  
>>> duplicate_profile = {"user_id": 212, "name": "Tommy"}  
>>> result = db.profiles.insert_one(new_profile) # This is fine.  
>>> result = db.profiles.insert_one(duplicate_profile)  
Traceback (most recent call last):  
DuplicateKeyError: E11000 duplicate key error index: test_database.profiles.$user_id_1 du
```

## See also

The MongoDB documentation on [indexes](#)

---

Copyright © MongoDB, Inc. 2008-present. MongoDB, Mongo, and the leaf logo are registered trademarks of

MongoDB, Inc

Made with Furo

